

PROPOSED PHASE ONE • DISCUSSION DRAFT

VICE & ClearAuth AI

Virtual Insurance Collection Employee and Prior-Authorization Agent — Vision and Scope Confirmation

Field	Detail
Prepared by	InfiniAI
Prepared for	Cascade AI
Purpose	Confirm alignment on the intended Phase One workflow, boundaries, functional requirements, and outcomes for VICE and ClearAuth AI, and serve as the basis for written client sign-off before build begins.
Status	Discussion draft — for review and confirmation, not yet approved

Purpose of this document

Relay Health AI proposes to deliver two connected intelligent agents on a single shared platform core: VICE, which pursues unpaid insurance claims after a service has been billed, and ClearAuth AI, which secures prior authorization before a procedure is ever scheduled. This document is intended to confirm whether the end-to-end workflow, boundaries, and requirements described below reflect the practice's intended scope for Phase One, and to give the practice a single reference it can review, redline, and sign off on before implementation begins.

Requested confirmation

- Does the practice agree that the VICE and ClearAuth AI workflows below are technically and operationally appropriate for Phase One?
- Should Phase One deliver both agents together on the shared platform core, or should one workflow (VICE or ClearAuth AI) go first, as recommended in Section 11?
- Which capabilities should be delivered in the initial implementation, and which should be deferred to a later phase?
- What assumptions, system limitations, dependencies, or changes to the practice's current process should be reflected before build begins?

1. Phase One objective

The objective is to automate the repeatable, rules-based, document- and call-intensive work of chasing unpaid claims and securing prior authorizations, while preserving clinical judgment and financial decision-making with the practice's own staff.

Automation should handle	Practice staff should retain
Claim and authorization intake; electronic status and eligibility checks; payer-portal submissions; IVR navigation and payer calls; reason classification; document retrieval and submission; follow-up scheduling; reconciliation against remittance; exception detection; notifications; and dashboard reporting.	Medical-necessity judgment, peer-to-peer review participation, approval of appeals and reprocessing requests above the configured policy threshold, auto-submit policy decisions, exception resolution, and any action outside the agreed decision tree.

Guiding principles

- Electronic channels (EDI 270/271/276/277/278) are tried first because they are fastest and cheapest; payer portals and APIs are tried second; a phone call is the last resort, used only when nothing electronic resolves the case.
- Staff see only exceptions — cases the waterfall could not resolve on its own.
- The agent may prepare and recommend actions; medical-necessity judgment and any action outside agreed policy thresholds always remain a human decision.
- The EHR/practice-management system and clearinghouse remain the systems of record; Relay Health AI's platform reads from and writes back to them — it does not replace them.
- Every electronic check, portal submission, and phone call is logged with a timestamp, source, and reference number for audit and payer-dispute purposes.
- Auto-submit thresholds, escalation routing, and payer-requirement overrides are configuration the practice's own admin controls — never a code change (Section 7).

2. Phase One workflow

Both agents follow the same underlying principle — electronic first, portal second, a phone call only when nothing else resolves it, and a human decision for anything that isn't administrative.

2A. VICE — claims-collection workflow

A. Daily A/R ingestion (“morning pull”)

1. Pull every unpaid or outstanding claim from the EHR/practice-management system and the clearinghouse each morning.
2. Normalize into a single work queue with claim age, dollar amount, payer, and current status.
3. De-duplicate against claims already in flight or already resolved.

B. Per-claim work selection and channel waterfall

1. Decide the next action per claim — wait/recheck vs. act now — from age, dollar value, payer behavior, and status; the highest-value and most at-risk claims (e.g. approaching timely-filing limits) surface first.

2. Check status electronically first via EDI 276/277. Interpret the response — received/pending/no-payment-issued, denial codes — and resolve or escalate.
3. If the electronic status is ambiguous, query the payer portal or API before spending a phone call.
4. If electronic and portal checks do not resolve the claim, call the payer: navigate the IVR, wait on hold, reach a live representative, identify as an automated representative of the provider **only where the payer-requirement matrix permits it**, and ask what is preventing payment.

C. Reason capture, classification, and resolution

1. Capture the representative's answer verbatim (transcript) and as structured data.
2. Classify the reason — missing document, medical-necessity denial, timely-filing denial, coordination of benefits / other insurance, coding issue, in process, or other — and route to the matching handler.
3. Missing document: retrieve the permitted copy from the EHR/document system, submit it through the correct electronic workflow, and log what was sent, to whom, and when.
4. Timely-filing denial: check the clearinghouse's original submission date. If filed on time, assemble the appeal packet and proof of timely filing; route to a human for approval, or auto-submit only where the practice's own policy threshold allows it (Section 4). If no valid appeal exists, flag as a staff exception.
5. Other resolvable reasons: request reprocessing on the same call where that resolves the claim.

D. Follow-up and recheck

1. Create time-based follow-ups (e.g. "recheck in 10 days if still unpaid") and re-enter the claim into the waterfall automatically at the follow-up date.

E. Reconciliation and reporting

1. Reconcile resolved claims against remittance (835/ERA) to confirm the payment was recovered.
2. Roll up claims reprocessed and dollars recovered into the VICE dashboard (Section 5).

F. Exception routing

1. Anything the waterfall cannot resolve — negotiations, unusual denials, no valid appeal — routes to the human exception queue with full context.

2B. ClearAuth AI — prior-authorization workflow

A. Order intake

1. Receive patient, insurance, CPT code, diagnosis, and physician order — via an EHR trigger or a one-click "Get Authorization" action.
2. Validate that the packet is complete enough to act on; incomplete packets route to an exception.

B. Determine whether authorization is required (electronic first)

1. Confirm active coverage via real-time eligibility (EDI 270/271).
2. Determine whether the plan requires authorization for this CPT code via payer API, clearinghouse, Availity, or EDI 278.
3. If **no authorization is required**, record the result, date, source, and reference/transaction number in the patient record, and stop — no portal, no call.

C. Portal submission (when authorization is required)

1. Assemble the clinical packet — order, office/visit notes, imaging history, and required supporting documentation.

2. Open the correct payer/authorization portal, enter patient and procedure details, upload the packet, and submit.
3. Track the submission reference and continue following the request to resolution.

D. Voice-call fallback

1. If the portal cannot resolve it, call the insurer directly, identifying as calling on behalf of the ordering physician regarding prior authorization.
2. Navigate the IVR, wait on hold, verify requirements with the representative, and obtain the authorization number or status.

E. Outcome capture and exceptions

1. Record the authorization number, validity window, and reference number; write results back to the EHR/patient record.
2. Track valid-through dates so an approval nearing expiry before the scheduled procedure date surfaces proactively.
3. Surface only exceptions to staff — e.g. “Clinical review required — missing six weeks of conservative-treatment documentation.” **The agent never makes an independent medical-necessity judgment;** that decision always goes to a licensed staff member.

3. Status and classification logic

3A. VICE — reason-for-non-payment classification

Reason class	Required system behavior
Missing document	Retrieve the permitted document from the EHR and submit through the correct electronic workflow; log recipient and timestamp.
Timely-filing denial	Check the original clearinghouse submission date; prepare an appeal for human approval (or auto-submit within policy threshold) if filed on time.
Medical-necessity denial	Route to a licensed staff member — never resolved by the agent.
Coordination of benefits / other insurance	Flag for staff verification of primary/secondary payer order.
Coding issue	Flag for staff/coding review; agent does not alter submitted codes.
In process	Schedule a follow-up recheck; no action required yet.
Other / unclassified	Route to the exception queue with the full transcript and reference number.

3B. VICE claim status indicators

Indicator	Meaning	Required system behavior
Resolved (electronic)	Claim paid or explained without a call	Logged and closed — no call needed
Escalated to call	Electronic/portal checks inconclusive	Agent places outbound call; recorded and transcribed
Follow-up scheduled	Payer requested missing information or cited in-process status	Document retrieved/submitted, or a recheck date is set
Appeal prepared	Timely-filing denial with proof of on-time filing	Routed for human approval, or auto-submitted within the configured policy threshold
Needs human attention	No valid appeal, unusual denial, or negotiation required	Routed to the exception queue with full context

3C. ClearAuth AI authorization status indicators

Indicator	Meaning	Required system behavior
No authorization required	Confirmed electronically	Logged with date, source, and reference number
Submitted — pending	Portal submission in progress	Tracked to resolution; escalate to a call if no response within the agreed window
Approved	Authorization number obtained	Auth #, valid-from/valid-to dates, and reference number written back to the EHR

Indicator	Meaning	Required system behavior
Clinical review required	Missing documentation or a medical-necessity question	Flagged to a licensed staff member; the agent does not decide
Expiring soon	Approval nearing its valid-to date before the scheduled procedure	Proactively surfaced to staff

4. Human-control boundaries and operating modes

The agent may	The agent may not without explicit practice authorization
Read and classify documents; run electronic and portal checks; navigate IVRs and speak with payer representatives; retrieve and submit permitted documents; monitor and update claim/authorization status; schedule follow-ups; reconcile against remittance; prepare appeal packets and exception reports; assemble reference and audit records.	Decide whether a claim or appeal should be pursued outside the practice's configured policy thresholds; make or imply a medical-necessity judgment; participate in a peer-to-peer review; auto-submit an appeal above the configured threshold; or alter historical records without an auditable, versioned change.

Operating-mode trust ramp

Autonomy is granted per action type, per workflow, and per payer — never all at once — and only as the exception rate proves the agent trustworthy.

Mode	Behavior	When enabled
Shadow	Agent proposes every action; a human approves each one before it goes out.	Recommended starting mode for the pilot (Section 11).
Supervised auto-submit	Low-risk actions (electronic checks, routine reprocessing requests) go live automatically; appeals and anything over the configured threshold still route to a person.	Enabled per workflow and per payer once shadow-mode accuracy is validated.
Wider autonomy	Broader auto-submit thresholds after a sustained low exception rate.	Enabled only by explicit, written practice sign-off, per workflow and per payer.

5. Expected Phase One outputs

- A daily-refreshed VICE dashboard: outstanding A/R, claims the agent is working, payers called, claims reprocessed, appeals submitted, payments recovered, and cases needing human attention.
- A ClearAuth AI outcome view per order: an approval (authorization number, valid dates, reference number) or a named exception flag.
- Reconciled claim and remittance status with reason classification and exception routing.
- Authorization results written back to the EHR/patient record with date, source, and reference number.
- A write-once audit and compliance log of every electronic check, portal submission, and call, with timestamp and reference number.
- Call recordings and transcripts retained per the practice's HIPAA retention policy.
- A reference/transaction-number ledger tying every claim or authorization reference back to the exact call, portal submission, or EDI transaction that produced it.
- An exception dashboard and approval queue, role-scoped to billing, front-desk, and clinical staff.

6. Systems and integration expectations

Every claim or authorization the agent touches passes through some combination of these integration points. None are optional — the agent is only as reliable as its weakest connection, so each one gets its own credentials, its own failure handling, and its own place in the audit log.

System	Expected role
EHR / practice-management system	Source of patient, insurance, CPT, order, and clinical-document data; destination for authorization and claim-status write-back.
Clearinghouse (e.g. Gateway EDI / TriZetto)	EDI 270/271/276/277/278 transactions — eligibility, claim status, authorization, remittance (835/ERA).
Payer portals & APIs	Availity, Optum, UHC, and other payer-specific portals for authorization submission and claim-status lookups.
Voice / telephone channel	IVR navigation and live conversation with payer representatives, for both claim-status and prior-authorization calls.
Appeals & denials workflow	Timely-filing checks, appeal-packet assembly, and the human-approval gate before submission.
Document & fax connector	Retrieval of operative reports and clinical packets from the EHR/document system; submission to payers where no portal exists.
Exception dashboard	The staff-facing review and approval queue, role-scoped by persona.
Audit & compliance log	Reference numbers, timestamps, and call recordings — write-once, retained per policy.

Connector types

Type	Description
API	Direct, versioned API — the fastest and most reliable option when a system offers one.
MCP	Model Context Protocol — structured, permissioned tool calls for agent actions, not just data reads.
Browser / portal automation	Drives a payer portal or legacy web interface directly when no API exists — the fallback, not the default.
Webhook / event	Real-time pushes from a client system instead of polling, where supported.
File / batch (SFTP)	Scheduled file drops — 835 remittance batches, nightly exports — for systems that only speak in files.

7. Platform, tenancy, and onboarding requirements

Deployment modes

Mode	Description
Multi-tenant SaaS (default)	Platform-hosted, shared infrastructure, strict per-tenant logical data isolation. Fastest to onboard, usage-metered.
Air-gapped / dedicated (on request)	Deployed inside the practice's own environment — on-premises or a private cloud — with no external network dependency. Same core and connector framework, chosen where compliance posture requires full isolation.

Three-tier admin model

Role	Responsibilities
Master Admin (platform operator)	Onboards new client tenants; manages the shared connector registry; sees cross-tenant usage and health metrics only — never PHI.
Tenant Admin (the practice's own admin)	Manages the tenant's users and roles; configures auto-submit thresholds, escalation routing, and payer-matrix overrides through the platform's rules engine; views the tenant's full audit log.
End User (practice staff)	Logs in to their own persona; sees only their role's queue and exceptions — nothing cross-tenant, nothing outside their permission.

Onboarding flow

1. The practice completes a guided questionnaire covering its EHR, clearinghouse, payer portals, and preferred deployment mode.
2. The platform checks its connector registry for a match; a known system means the tenant is auto-provisioned in minutes.
3. An unmatched system routes to the platform admin to scope and build a new connector, which then joins the registry for all future clients.
4. The practice supplies credentials through a secure vault, scoped per connector.
5. Each connection is validated with a test call or API ping before go-live.
6. The tenant goes live and users are invited to log in under their own persona.

Discovery module

Before build begins, Relay Health AI conducts structured working sessions with the staff who currently do this work — billing, prior-authorization coordinators, front desk — covering systems and ownership, claim/authorization volume, current-state process, rule ownership, MVP scope, and data validation (25 questions across six categories; see Appendix A). The output is a scoped MVP and build plan grounded in the practice's actual numbers, not a template. Discovery also establishes a shadow/comparison window — the agent's proposed action checked against what staff actually did — before the agent runs unsupervised.

8. Compliance and guardrails

This is the boundary the whole design is built around: administrative and procedural work is fully automatable; medical judgment is not. Every guardrail below exists to keep that line firm and provable.

Guardrail	Requirement
HIPAA / BAA enforcement	A signed Business Associate Agreement is required with every subprocessor in the chain — the voice platform, hosting provider, and any model API — before PHI touches the system.
Call disclosure control	The agent identifies itself as an automated representative only where the payer permits it; the payer-requirement matrix tracks which payers accept this versus require a human caller.
Recording-consent handling	One-party vs. two-party consent rules are checked per call by state and by payer — never assumed.
No-medical-judgment boundary	Medical necessity, peer-to-peer review, and unusual denials always escalate to a licensed staff member; the agent reports facts and missing documentation only.
Data handling	PHI is encrypted at rest and in transit; retention follows the practice's existing HIPAA policy rather than a separate schedule.
Client IP & ownership	Every case record, configured rule, and document the practice creates belongs to the practice — exportable on request, and never used to train a model shared across other tenants without explicit opt-in.
Configurable, audited rules	Auto-submit thresholds and escalation routing are set by the Tenant Admin, never hardcoded; every rule change is itself an audited, versioned event.
Tenant isolation	Logical isolation per tenant in the multi-tenant SaaS deployment; physical isolation in the air-gapped deployment — a breach or bug in one tenant cannot cross into another's.

9. Operational knowledge to establish and maintain

The integrations in Section 6 are the pipes. What keeps the agent accurate day to day is a set of living references someone has to own, update, and audit — otherwise the agent starts making calls with stale information.

Static knowledge — changes slowly, reviewed on a schedule

- Payer-requirement matrix — which payers require prior authorization for which CPT codes, and which permit an automated caller; the single biggest source of stale-data errors if left unreviewed.
- IVR call maps and representative scripts, per payer — the phone trees, hold patterns, and phrasing that gets a straight answer, captured once and reused every call.
- Auto-action policy thresholds — the practice's own rules for what the agent may submit or appeal automatically versus what always needs a signature.

Living operational state — changes every day

- Credential and access vault — clearinghouse, payer-portal, and EHR credentials, scoped per connector, rotated on schedule, never shared across integrations.
- Escalation runbook — which staff member handles clinical-necessity questions, peer-to-peer review calls, and denials that don't fit the standard decision tree.

Audit trail — write-once, kept for dispute resolution

- Reference/transaction-number ledger — every claim or authorization reference number tied back to the exact call, portal submission, or EDI transaction that produced it.
- Call recordings and transcripts — kept against the practice's retention policy, the only way to resolve a dispute over what a payer representative actually said.

10. What we need from you

None of this is unusual for a practice already running an EHR and a clearinghouse — it is mostly granting the agent the same access a new staff member would need, plus a few decisions only the practice can make.

Access and credentials

- EHR / practice-management integration access — read access to patient, insurance, claim, and order data; write access to log authorization results and attach documents.
- Clearinghouse credentials for eligibility, claim-status, and authorization transactions.
- Payer-portal logins — Availity and any payer-specific portals already in use for authorizations and claim status.
- A dedicated outbound phone line, separate from the practice's main line, with call recording enabled.

Policy decisions only the practice can make

- Auto-submit thresholds — which appeals and reprocessing requests can go out automatically, and which always wait for a signature.
- Pilot scope — which payers and which claim/procedure types are in scope for the first phase; narrower is better for the first pilot.

- Exception owner — one named staff member who triages the exception queue during the pilot.

Legal and data

- A signed Business Associate Agreement covering every vendor in the chain that touches PHI — the voice platform, hosting, and any model API.
- A sample of unpaid claims and past authorization outcomes, used to validate the decision logic before it runs live.

11. Engagement model and phased rollout

Engagement stages

Stage	What it includes
1 • Discovery & onboarding	Working sessions with current-state staff; full systems, payer, and volume inventory; MVP definition and a phase-2/3 forecast; data-validation sample review; a written discovery report and scoped proposal.
2 • Build the dedicated solution	EHR/PM, clearinghouse, and payer-portal integration work; voice-agent conversation design tuned to each in-scope payer's IVR; orchestration/waterfall logic and the practice's auto-submit policy rules; exception dashboard and audit-log build-out; HIPAA safeguards and BAA setup with every subprocessor.
3 • Run in production	Cloud hosting, voice-platform usage, telephony connectivity, clearinghouse/EDI transaction fees, audit-log storage and monitoring, and ongoing payer-rule-matrix maintenance — metered against actual use.

Recommended pilot phasing

The recommended first pilot is one workflow, one payer set, and a shadow period before anything submits automatically — so the practice can see the agent's decisions before trusting it to act on them unsupervised.

Phase	Description
Phase 1 — Pick one workflow	Either VICE (claims follow-up) or ClearAuth AI (prior authorization) — not both at once — scoped to the practice's highest-volume payer.
Phase 2 — Shadow mode	The agent runs the full waterfall and proposes every action, but a human approves each one before it goes out; validates accuracy against real cases.
Phase 3 — Supervised auto-submit	Low-risk actions go live automatically; appeals and anything over the practice's threshold still route to a person.
Phase 4 — Widen scope	Add payers, add the second workflow, and raise auto-submit thresholds only as the exception rate proves the agent is trustworthy.

12. Confirmation requested from you

1. Does this workflow match your vision for VICE and ClearAuth AI, as described in Sections 1–5?
2. Should Phase One deliver both agents together, or should one workflow go first, scoped to your highest-volume payer (Section 11)?
3. Which steps should be automated through APIs/EDI, deterministic rules, browser/portal automation, or a voice call, given your current systems?
4. What EHR/practice-management custom fields or process changes are needed for authorization status, claim-reason codes, reference-number storage, and document completeness?
5. How should duplicate claim/authorization creation, incorrect document matching, and unauthorized external notifications be prevented, beyond what's described in Section 4?
6. What human-approval gates, audit-trail retention period, and confidence thresholds should be in place before wider autonomy (Section 4) is granted?

7. What measurable acceptance criteria should apply to extraction accuracy, classification accuracy, call-outcome accuracy, and system reliability during the pilot?
8. What elements, if any, should be deferred to a subsequent phase, and why?
9. Which deployment mode — multi-tenant SaaS or dedicated/air-gapped — fits your compliance posture (Section 7)?

13. Closing

Relay Health AI views this as a practical first deployment because it replaces two of the most repetitive, phone-heavy administrative burdens in a physician practice or outpatient rehab center's revenue cycle, while keeping every medical-necessity decision, appeal threshold, and payment action under the practice's own control. We welcome your feedback on the proposed workflow, the required process changes, and whether this scope reflects your intended definition of Phase One.

Sign-off

By signing below, the parties confirm agreement with the Phase One vision, scope, and requirements described in this document as the basis for the Discovery stage in Section 11.

Role	Name / title	Signature	Date
Prepared by — InfiniAI			
Reviewed by — Relay Health AI			
Approved by — Practice / Client			
Approved by — Practice Compliance / Privacy Officer (if applicable)			

14. Appendix A — Discovery question bank

The 25 questions below, grouped into six categories, are the working-session agenda used during Discovery (Section 7) to ground every payer-requirement matrix, IVR script, and auto-submit threshold in the practice's actual current-state process rather than a template.

A • Systems & ownership — what you run today

- 1. Which EHR / practice-management system do you use, and who administers it day to day?
- 2. Which clearinghouse do you use, and which transaction types (eligibility, claim status, authorization) are actually turned on versus available but unused?
- 3. Which payer portals does staff log into today, and who holds those logins?
- 4. What phone system do you use for payer calls, and are calls already recorded?
- 5. Is there a spreadsheet, whiteboard, or tracker staff actually use instead of the official system — and can we see it?

B • Volume — what we're sizing this for

- 1. What is your total outstanding A/R today, broken down by 0-30 / 31-60 / 61-90 / 90+ days?
- 2. How many new claims do you file per month, and how many prior-authorization requests per month?
- 3. Which 5-10 payers cover most of your claim count and dollar volume?
- 4. Which CPT / procedure codes trigger prior authorization most often, and how frequently does each come up?
- 5. Roughly how many payer phone calls does staff place per week — separately for claims follow-up and for prior auth?

6. How much time does staff spend, on average, per claim and per authorization today?

C • Current-state process — how it's actually done

1. Can you walk us through one real claim, start to finish, with the person who works it?
2. Can you walk us through one real prior-authorization request the same way?
3. What makes staff decide to call the payer right now versus wait and check again later?
4. What gets escalated to a supervisor or physician today, and through what channel?
5. Where do exceptions and follow-ups get tracked today, if anywhere?

D • Rule ownership — who actually knows the policy

1. Who maintains your payer-requirement matrix, and how is it kept updated?
2. Who decides what can be submitted or appealed automatically versus what always needs a signature?
3. Who is the clinical escalation contact for medical-necessity questions and peer-to-peer reviews?
4. How does staff find out when a payer changes its requirements today?

E • MVP scope vs. future forecast

1. What's the smallest slice worth automating first — one payer, one CPT family, or claim-status calls only?
2. What claim/authorization volume would make a pilot feel real, without being unmanageable to validate by hand?
3. What should be explicitly out of scope for v1, and deferred to a later phase?

F • Data validation

1. Can you provide a sample export of real claims and authorization history for us to test the decision logic against?
2. Can we agree on a comparison window — the agent's proposed action checked against what staff actually did — before it runs unsupervised?